

Holding BERTs Hands Through Multiple Choices

Wumbolo, Tegla, Nikola Kraljević

University of Zagreb, Faculty of Electrical Engineering and Computing

Unska 3, 10000 Zagreb, Croatia

{Wumbolo, Tegla, Nikola.Kraljevic2}@fer.hr

Abstract

Question answering tasks have long been a central challenge within NLP. This paper presents our approach to tackling the Sentence Puzzle task in the SemEval-2024 Task 9 competition: "BRAIN-TEASER: A Novel Task Defying Common Sense." Diverging from the conventional multi-class classification framework, we introduce a novel pairwise comparison methodology. Instead of forcing the model to choose the correct answer from four options, we reformulate the task into multiple decisions, evaluating triplets of answers at a time. This approach, inspired by the efficiency of One-vs-One and One-vs-Rest strategies, aims to enhance the model's reasoning capabilities and accuracy. Our results showcase... ako budu dobri rezultati tu cemo flexat, xD

1. Introduction

The competition, named "SemEval 2024 BRAINTEASER: A Novel Task Defying Common Sense" described in the paper (Jiang et al., 2023) is here to push modern models to their limits, putting them to a quite adversarial setting. The paper makes a distinction between vertical and lateral thinking. Vertical thinking, also known as linear, logical, or convergent, is mostly a sequential and analytical process, requiring direct memory recall or a few logical steps to come to a sensible conclusion. Lateral thinking, also known as "thinking outside the box," is more of a divergent and creative process, where the question might not make sense when reading it at first, and to come up with an answer, you need to explore multiple angles. While LLMs show good vertical thinking capabilities; they are notorious for hallucinating answers. An even more adversarial task for LLMs is answering questions that require lateral thinking. State-of-the-art models such as ChatGPT show an accuracy of 60%, humans show a 90% accuracy, while random guessing gets close to 25% accuracy as there are four answer candidates for each question

In this paper, we tackle the "BRAINTEASER" puzzles with an alternative approach which we believe will make it easier for transformer models to reason about all the possible answers, focusing on BERT like models. In Section 3. we will describe the form of the two types of puzzles available. Our approach and the reasoning behind it will be described in Section 4. Results of our experiment will be presented in Section 5. and Section 6. respectively.

2. Related Work

Reasoning in NLP is a hot topic, with LLMs in the center of attention. It is no secret that LLMs are far from perfect, and it takes only one session of asking questions about a topic one is well versed in to see inconsistencies, but also be amazed at times. Knowing where a model fails is important and leads to improvement. Many benchmarks have been made with the purpose of challenging models with more intricate reasoning, like "Commonsense QA" (Talmor et al., 2019), a dataset consisting of questions that are easy for humans and require no prior knowledge, like a

specific document or context, just common sense. Similarly, the "BRAINTEASER" (Jiang et al., 2023) is a benchmark dataset consisting of questions which are constructed in such a way that it is required to consider multiple approaches when answering them.

As this was a competition dataset, there were some previous works that tackled this problem. A team of researchers from the National Technical University of Athens published their submission (Panagiotopoulos et al., 2023) named "Transformer Models for Lateral Thinking Puzzles". They showed promising results, significantly outperforming baselines of 60% reported in (Jiang et al., 2023). Their approach consisted of fine tuning models with a straightforward approach, treating the problem as a multi-class classification task. Additionally, they performed a transformation of the problem to a binary classification problem. The transformation took form of taking each question with four candidate answers and transforming it to three questions which required a binary label, signaling if the answer was correct or not ¹. Their results showed that the transformation was not quite useful, as the same models had somewhat worse accuracies when the transformation was applied.

This inspired us to explore another transformation of the problem, as we believe that subjecting a model to only one candidate answer takes away information. When one is answering a question with candidate answers, it is a good idea to eliminate candidates that are somewhat obviously wrong. That is why we believe an approach where we force a model to choose between two candidate answers at a time might be helpful, this is described in detail in Section 4.

3. Dataset

As we've mentioned before the "BRAINTEASER" dataset consists of two types of puzzles, "Sentence puzzles" and "Word puzzles". Word puzzles focus more on letter composition, defying the original meaning of the words at hand, shown in Table 1. The second type of puzzles we find more interesting, so called "Sentence puzzles", shown in Table 2. The three step process of constructing the "BRAIN-

¹The fourth candidate answer was always "None of the above" so they discarded it.

Table 1: A **word puzzle** question(sub-task B). The correct answer is in **bold**.

Question	Candidates
What part of London is in France?	The letter N. The letter O. The letter L. None of the above.

Table 2: A **sentence puzzle** examples(sub-task A). The correct answer is in **bold**.

Question	Candidates
A man shaves everyday, yet keeps his beard long.	He is a barber. He wants to maintain his appearance. He wants his girlfriend to buy him a razor. None of the above.

TEASER” dataset is described in great detail in (Jiang et al., 2023). The process consists of data collection (web scraping), distractor sampling and generating reconstruction examples. As some pre-trained models could have these questions in their training corpus, steps were made to ensure that the questions were novel for the models. The data set consists of original data, semantic reconstruction data, context reconstruction data. The semantic reconstruction used an open-source rephrasing² tool and human annotators as quality control to rephrase the questions. The context reconstruction process used GPT-4 to initially shift the context of the question, as well as human annotators. In the following sections we will refer to these distinct types of questions as **Ori.**, **Sem.** and **Con.** for original questions and semantic/context reconstruction respectively. Lastly, the data was split into 507/120/120 questions for the train/test/validation datasets respectively.

4. Pairwise approach

We perform a transformation of the questions in the following way. Each tuple of a question and its candidate answers

$$(Q, A, B, C, D)$$

gets transformed into three tuples:

$$\begin{aligned} &(Q, A, B, D), \\ &(Q, A, C, D), \\ &(Q, B, C, D). \end{aligned}$$

In case the correct answer was C , then the ground truth label in the tuple (Q, A, B, D) is D (None of the above).

5. Results

Eval Results

6. Conclusion

Conclusion

References

- Yifan Jiang, Filip Ilievski, Kaixin Ma, and Zhivar Sourati. 2023. Brainteaser: Lateral thinking puzzles for large language models.
- Ioannis Panagiotopoulos, Giorgos Filandrianos, Maria Lymperaio, and Giorgos Stamou. 2023. Ails-ntua at semeval-2024 task 9: Cracking brain teasers: Transformer models for lateral thinking puzzles.
- Alon Talmor, Jonathan Herzig, Nicholas Lourie, and Jonathan Berant. 2019. Commonsenseqa: A question answering challenge targeting commonsense knowledge.

²<https://quillbot.com/>

Table 3: Sentence puzzle

System	Original	Semantic	Context	Ori. + Sem.	Ori. + Sem. + Con.	Overall
Multi-class classification problem						
Human	.907	.907	.944	.907	.889	.920
ChatGPT	.608	.593	.679	.507	.397	.627
RoBERTa-L	.435	.402	.464	.330	.201	.434
microsoft/deberta-v3-base	.775	.775	.700	.775	.675	.750
FacebookAI/roberta-large	.850	.850	.775	.850	.700	.825
Binary classification problem						
microsoft/deberta-v3-base	.650	.650	.525	.625	.625	.608
FacebookAI/roberta-large	.125	.125	.125	.125	.125	.125
Pairwise						
microsoft/deberta-v3-base	.650	.570	.650	.650	.525	.623
Facebook/roberta-large	.675	.700	.725	.650	.600	.700

Table 4: Word puzzle

System	Original	Semantic	Context	Ori. + Sem.	Ori. + Sem. + Con.	Overall
Multi-class classification problem						
Human	.917	.917	.917	.917	.900	.917
ChatGPT	.561	.524	.518	.439	.292	.535
RoBERTa-L	.195	.195	.232	.146	.061	.207
FacebookAI/deberta-v3-base	.625	.688	.625	.594	.375	.646
FacebookAI/roberta-large	.312	.375	.438	.250	.188	.375
Binary classification problem						
microsoft/deberta-v3-base	.000	.000	.000	.000	.000	.000
FacebookAI/roberta-large	.000	.000	.000	.000	.000	.000
Pairwise						
microsoft/deberta-v3-base	.562	.469	.500	.469	.250	.510